

---

# Course Project: Part 1

Generative AI

Saarland University – Winter Semester 2024/25

---

Martínez

7057573

cama00005@stud.uni-saarland.de

## 1 Part 1.a: Basic Prompts

- 2 (I.1) Modify the parameters in the `project_part1_evaluate.py` script to query GPT-4o-mini
- 3 with the basic prompt to generate program repairs and run the script. Report *RPass* and
- 4 *REdit*.
- 5 (I.2) Repeat the steps in (I.1) for Phi-3-mini and report your results.

Table 1: Overall results for GPT-4o-mini and Phi-3-mini with the basic prompt to generate program repairs. For (I.1) and (I.2).

Model	<i>RPass</i>	<i>REdit</i>
GPT-4o-mini	88.0	23.77
Phi-3-mini	36.0	18.11

Table 2: Average per problem results for GPT-4o-mini with the basic prompt to generate program repairs. The average is calculated without considering the cases with edit distance = −1. For (I.1).

Problem	Prog. 1	Prog. 2	Prog. 3	Prog. 4	Prog. 5	Avg.
Problem 1	4	−1	15	19	24	15.50
Problem 2	1	63	16	85	54	43.80
Problem 3	7	19	11	36	10	16.60
Problem 4	53	5	5	39	2	20.80
Problem 5	−1	12	4	−1	39	18.33

Table 3: Average per problem results for Phi-3-mini with the basic prompt to generate program repairs. The average is calculated without considering the cases with edit distance = −1. For (I.2).

Problem	Prog. 1	Prog. 2	Prog. 3	Prog. 4	Prog. 5	Avg.
Problem 1	23	−1	5	−1	−1	14.00
Problem 2	−1	−1	−1	−1	−1	0.00
Problem 3	−1	−1	−1	−1	10	10.00
Problem 4	51	4	5	−1	30	22.50
Problem 5	−1	−1	10	−1	25	17.50

- 6 (I.3) Modify the parameters in the `project_part1_evaluate.py` script to query GPT-4o-mini  
7 with the basic prompt to generate hints and run the script. Evaluate the generated hints  
8 using the *HGood* metric [1] and report your results.

Table 4: Hint Quality Metrics for GPT-4o-mini.

Problem	Program	<i>HCorrect</i>	<i>HInformative</i>	<i>HConceal</i>	<i>HComprehensible</i>	<i>HGood</i>
Problem 1	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 1		1.0	1.0	1.0	1.0	1.0
Problem 2	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 2		1.0	1.0	1.0	1.0	1.0
Problem 3	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 3		1.0	1.0	1.0	1.0	1.0
Problem 4	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 4		1.0	1.0	1.0	1.0	1.0
Problem 5	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 5		1.0	1.0	1.0	1.0	1.0
Overall Average		1.0	1.0	1.0	1.0	1.0

- 9 (I.4) Repeat the steps in (I.3) for Phi-3-mini and report your results.

Table 5: Hint Quality Metrics for Phi-3-mini.

<b>Problem</b>	<b>Program</b>	<b><i>HCorrect</i></b>	<b><i>HInformative</i></b>	<b><i>HConceal</i></b>	<b><i>HComprehensible</i></b>	<b><i>HGood</i></b>
Problem 1	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 1		1.0	1.0	1.0	1.0	1.0
Problem 2	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	0	0
	Prog. 4	0	1	1	1	0
	Prog. 5	1	1	1	1	1
Avg. Problem 2		0.8	1.0	1.0	0.8	0.6
Problem 3	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	0	0
	Prog. 4	1	1	1	0	0
	Prog. 5	1	1	1	1	1
Avg. Problem 3		1.0	1.0	1.0	0.6	0.6
Problem 4	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	0	0
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	0	0
Avg. Problem 4		1.0	1.0	1.0	0.6	0.6
Problem 5	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	0	0
	Prog. 3	0	1	1	0	0
	Prog. 4	1	1	1	0	0
	Prog. 5	0	0	1	0	0
Avg. Problem 5		0.6	0.8	1.0	0.2	0.2
Overall Average		0.88	0.96	1.0	0.64	0.60

## 10 Part 1.b: Improved Program Repairs

11 (I.5) Modify the scripts to query GPT-4o-mini for generating  $n = 3$  repair candidates with a  
 12 temperature of 0.7 and select the best candidate. Next, run the scripts and report  $R_{Pass}$  and  
 13  $R_{Edit}$ .

14 The modifications I did to the provided codebase to implement the *improved program repairs* are as  
 15 follows:

- 16 1. The method `Repair.generate_repair()` now takes 3 additional positional pa-  
 17 rameters: `n=1`, `do_sample=None`, `temperature=None`. These are defined in  
 18 the main method of the `project_part1_evaluate.py` and then propagated  
 19 through `DatasetEvaluation.evaluate_programs_in_folder()`, which calls the  
 20 method `ProgramEvaluation.get_and_evaluate_repair()`, which in turn calls the  
 21 `Repair.generate_repair()`, defined in `project_part1_repair.py`. This allows to  
 22 turn on/off sampling and set the temperature for the generation of repair candidates, without  
 23 affecting the basic workflow of **Part 1.a**.
- 24 2. The temperature parameter, which was propagated through as explained above, is now  
 25 accepted as a new parameter in the `Repair.call_llm_openai()` function. Namely, when  
 26 querying GPT-4o-mini and Phi-3-mini, it is set to 0.7.
- 27 3. Similarly, the `Repair.call_llm_huggingface()` now receives these 2 additional param-  
 28 eters: `do_sample` and `temperature`, which are then passed to `self.model.generate()`  
 29 as additional positional arguments. This allows to enable sampling and set the temperature,  
 30 when querying Phi-3-mini.
- 31 4. Inside `Repair.generate_repair()`, the algorithm 1 was implemented.

---

**Algorithm 1:** Algorithm for *improved program repairs*.

---

**Input :** Problem data  $T$ , buggy program  $P_b$ , test suite  $\Omega_T$ , sampling flag  $s$ , temperature  $\tau$ ,  
 number of repaired programs to query  $n$

```

1 generateRepair ( $T, P_b, \Omega_T, n, s, \tau$ )
2    $\mathcal{C} \leftarrow \text{COMPILER}()$                                 // Compiler instantiation
3    $\mathcal{D} \leftarrow \text{DISTANCE}()$                                 // Distance metric initialization
4    $\mathcal{S} \leftarrow \emptyset$                                     // Initialize repair statistics set
5   Repeat
6      $r \leftarrow \text{QUERYLLM}(T, P_b, s, \tau)$                 // Generate repair candidate
7      $f \leftarrow \text{EXTRACTFIX}(r)$                             // Extract fixed program
8
9     // The result of the test suite are stored in a list of tuples  $(\omega, o)$  where
10    //  $\omega$  represents whether the test case passed or not, and  $o$  is the output
11    // of the test case
12     $\Omega_{\text{results}} \leftarrow \mathcal{C}(f, \Omega_T)$                 // Execute test suite  $\Omega_T$  on  $f$ 
13
14    // Count the number of passing test cases
15     $\omega_c \leftarrow \sum_{(\omega, o) \in \Omega_{\text{results}}} \mathbf{1}(\omega)$         //  $\mathbf{1}(\cdot)$ : Indicator function for test success
16     $\delta \leftarrow \mathcal{D}(f, P_b)$                             // Compute edit distance between  $f$  and  $P_b$ 
17     $\mathcal{S} \leftarrow \mathcal{S} \cup \{(f, \omega_c, \delta, \Omega_{\text{results}})\}$     // Store the statistics of the current repair
18  until  $n$  times
19
20  // Total ordering: Maximize  $\omega_c$ , minimize  $\delta$ . If there are multiple correct  $f$ ,
21  // choose the one with smallest  $\delta$ 
22   $\mathcal{S}_{\text{max}} \leftarrow \{(f, \omega, \delta, \Omega) \in \mathcal{S} \mid \omega = \max_{(f', \omega', \delta', \Omega') \in \mathcal{S}} \omega'\}$ 
23   $(P_r, -, \delta_{\text{opt}}, \Omega_{\text{results}}) \leftarrow \arg \min_{(f, \omega, \delta, \Omega) \in \mathcal{S}_{\text{max}}} \delta$ 
24  return  $(P_r, \Omega_{\text{results}})$ 

```

**Output :** The repaired program  $P_r$ , with the most correct testcases and closest to  $P_b$ , and the  
 associated results of evaluating it with the testcases,  $\Omega_{\text{results}}$ .

---

- 
- 33 (I.6) Repeat the steps in (I.5) for Phi-3-mini and report your results.

Table 6: Overall results for GPT-4o-mini and Phi-3-mini with the basic prompt to generate program repairs. For (I.5) and (I.6).

<b>Model</b>	<b><i>RPass</i></b>	<b><i>REdit</i></b>
GPT-4o-mini	96.0	22.71
Phi-3-mini	60.0	25.87

34 **Part 1.c: Advanced Workflow**

- 35 (I.7) Modify the required scripts to improve the quality of the hint generation process of GPT-4o-  
 36 mini. Evaluate the generated hints using the *HGood* metric [1] and report your results.

Table 7: Hint Quality Metrics for GPT-4o-mini with the *advanced workflow*.

Problem	Program	<i>HCorrect</i>	<i>HInformative</i>	<i>HConceal</i>	<i>HComprehensible</i>	<i>HGood</i>
Problem 1	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 1		1.0	1.0	1.0	1.0	1.0
Problem 2	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 2		1.0	1.0	1.0	1.0	1.0
Problem 3	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 3		1.0	1.0	1.0	1.0	1.0
Problem 4	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 4		1.0	1.0	1.0	1.0	1.0
Problem 5	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 5		1.0	1.0	1.0	1.0	1.0
Overall Average		1.0	1.0	1.0	1.0	1.0

37 (I.8) Repeat the steps in (I.7) for Phi-3-mini and report your results.

Table 8: Hint Quality Metrics for Phi-3-mini with the *advanced workflow*. The CoT prompting seems to have addressed the formatting and readability issues but at the cost of some specificity in the hints, as the *HInformative* metric is lower than in the basic prompt approach.

Problem	Program	<i>HCorrect</i>	<i>HInformative</i>	<i>HConceal</i>	<i>HComprehensible</i>	<i>HGood</i>
Problem 1	Prog. 1	1	0	1	1	0
	Prog. 2	1	1	1	1	1
	Prog. 3	1	0	1	1	0
	Prog. 4	1	0	1	1	0
	Prog. 5	1	1	1	1	1
Avg. Problem 1		1.0	0.4	1.0	1.0	0.4
Problem 2	Prog. 1	1	0	1	1	0
	Prog. 2	1	0	1	1	0
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 2		1.0	0.6	1.0	1.0	0.6
Problem 3	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	0	1	1	0
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 3		1.0	0.8	1.0	1.0	0.8
Problem 4	Prog. 1	0	0	1	1	0
	Prog. 2	1	1	1	1	1
	Prog. 3	1	1	1	1	1
	Prog. 4	1	1	1	1	1
	Prog. 5	1	0	1	1	0
Avg. Problem 4		0.8	0.6	1.0	1.0	0.6
Problem 5	Prog. 1	1	1	1	1	1
	Prog. 2	1	1	1	1	1
	Prog. 3	1	0	1	1	0
	Prog. 4	1	1	1	1	1
	Prog. 5	1	1	1	1	1
Avg. Problem 5		1.0	0.8	1.0	1.0	0.8
Overall Average		0.96	0.64	1.0	1.0	0.64

38 (I.9) Explain the approach you took to improve the results of the models in detail. Mention some  
39 of the limitations of this approach and ideas of how it could be further improved in future  
40 works.

41 The approach I took to create the *advanced workflow* and improve the results of the models involved  
42 making the following modifications to the provided codebase:

- 43 1. A new parameter called `use_advanced_workflow` was introduced in the `Hint` class. This  
44 allows to turn on the **Part 1.c**'s *advanced workflow* by setting it to `True`, and the basic  
45 workflow used for **Part 1.a** and **Part 1.b** by setting it to `False`.
- 46 2. The function `Hint.generate_hint()` was modified to include a new positional ar-  
47 gument, `repair_agent`. So that, if `use_advanced_workflow` is set to `True`, the  
48 `repair_agent.generate_repair()` function is called, obtaining the repaired program  
49  $P_r$  and the results of the test cases  $\Omega_{\text{results}}$ , as outlined by the algorithm 1.
- 50 3. The parameters `repaired_program` and `testcases_results` were added to the func-  
51 tion `Hint.call_llm()`, so that `Hint.generate_hint()` can pass them to generate the  
52 hint, upon calling the latter. Moreover, if `use_advanced_workflow` is set to `True`, the  
53 `user_prompt` is formatted to include not only the problem data and the student's buggy  
54 program, but also the repaired program  $P_r$  and the results of the test cases  $\Omega_{\text{results}}$ . See  
55 **Appendix: Advanced Prompt**, which shows the template of the advanced prompt used.

56 This approach allowed me to improve the basic prompt by adding the repaired program  $P_r$  to the  
57 model, as well as the evaluation of  $P_r$  on the test suite  $\Omega_T$ . This provides more context to the model,  
58 whilst taking into account that, since  $P_r$  is also generated by the model and thus is not guaranteed to  
59 be correct, the model *knows* how correct  $P_r$  actually is, based on the results of the test cases. This  
60 allows the model to have a better understanding of the problem, how a potential solution should look  
61 like (the provided repaired program), how good the potential solution is (the testcases results), and  
62 uses this context to gain ideas on potential things to pinpoint at, in order to create the final hint to the  
63 student.

64 This approach created a 2-step look-ahead process for the model to generate hints, which can be  
65 summarized as follows:

- 66 1. The model is prompted with a buggy code  $P_b$  to generate a repair, based on the problem  
67 data  $T$ .
- 68 2. Using **Part 1.b**'s *improved program repairs*, it generates a repaired program  $P_r$  and evaluates  
69 it on the test suite  $\Omega_T$ .
- 70 3. The repaired program  $P_r$  is added to a new prompt, which also includes the problem data  
71  $T$ , the buggy code  $P_b$ , and the results of the test cases  $\Omega_{\text{results}}$ . This new prompt is used to  
72 generate the hint.
- 73 4. Upon prompting the same model with this new prompt, we give the chance for the model to  
74 analyze the problem again, the student's buggy code and a potential solution  $P_r$ , which from  
75 the model's perspective should be correct, as it was prompted in the first step to specifically  
76 generated a *repaired* program given  $P_b$  and  $T$ . But now that we included the testcases results  
77 inside the given prompt, the model can see how wrong or correct the produced  $P_r$  was, and  
78 we give it the opportunity to rationalize on why the  $P_r$  that it produced on the first step was  
79 wrong.
- 80 5. The model has now seen the problem twice. It has also seen a  $P_r$ , which was produced by  
81 itself in the first step, and it knows how correct or incorrect it was, based on the testcases  
82 results. With this information and the constrain to make it create a Chain of Thought inside  
83 the [EXP] placeholders, the model can focus not only on  $P_b$ , but also on how a potential  
84 solution  $P_r$ , that a student could have made, may also be wrong and why. In the case that  
85  $P_r$  is correct, the model can directly focus on providing a hint to the student that makes  $P_b$   
86 closer to the  $P_r$ , that it now knows for certain that is correct, thanks to the testcases results.

87 The previous explanation is cleverly illustrated in the Figure 5, in which the Phi-3-mini model is  
88 not able to generate a correct repair, but given that it knows that it is not correct, the model in its  
89 Chain-of-Thought explanation pinpoints bugs in the repaired program, as well as in the student's  
90 buggy code. This steers the model off of potentially providing the student a hint that would make it



91 go towards the wrong repaired program (since in the first step, when it was prompted to generate a  
92 repair, it thought that this particular repaired program was correct). This can be seen when the model  
93 says: "(...) *The problem with the repaired code is that it does not maintain the relative order of the*  
94 *elements in the original list, which is a key requirement of the problem statement (...)*". The model  
95 now has to *think*<sup>1</sup> twice and figure out how the student's buggy code and the repaired program are  
96 wrong, and thus it can provide a hint that looks ahead on problems that the student may encounter.  
97 This can be seen in the hint that the model finally provided: "(...) *Consider using a data structure that*  
98 *maintains the order of elements as you iterate through the list and check for duplicates. (...)*". There,  
99 it took into account both problems, the one in the repaired program as well as one in the student's  
100 buggy code, and provided a hint that combined both. In the case of GPT-4o-mini, since it was already  
101 quite good at generating hints, as seen in Table 5, the effect of this approach was not as pronounced,  
102 but it still allowed the model to provide even better hints.

103 **Limitations.** There are of course some limitations to this approach. The better the model is in  
104 generating repairs, the better the hints will be. This can be seen on Table 7 and Table 8, in which it is  
105 clear that GPT-4o-mini is better in terms of the *HGood* metric compared to Phi-3-mini. If a model  
106 always provides correct repairs, the model does not have to analyze too much, as it can simply give a  
107 concise hint that directly makes the student's buggy code closer to the repaired program.

108 The model is also limited by the quality of the test cases, as it can only evaluate the repaired program  
109 based on the test cases results. If the test cases are not good, the model may generate a repaired  
110 program that is actually wrong, but the test cases may not be able to capture this, and thus the model  
111 may generate a hint that makes the student's buggy code closer to this wrong repaired program.

112 **Further Improvements**<sup>2</sup>. The approach could be further improved by adding even more context to  
113 the prompt. For example, the actual input and expected output of the testcases could be provided,  
114 instead of only saying whether the testcases passed or not. This would provide the model exactly  
115 what it's been evaluated on, and based on the problem data, it should deduce some logic out of the  
116 given testcases.

117 At the same time, few-shot learning could be used by providing the model with examples of problems,  
118 their corresponding buggy code, repaired code and testcases results, and a final hint that was given  
119 to the student. This hint would have already been externally evaluated and deemed as good. This  
120 would allow the model to learn from examples of good hints, learn what actually makes a good  
121 hint, and generate better hints in the future that resemble those examples. In practice and if it is not  
122 possible for a human to review all examples and hints, GPT-4o-mini or an even better (potentially  
123 larger and non-free) model could be used to generate these examples with some fixed budget, and  
124 then Phi-3-mini or another open-source model could be used in production to generate hints based on  
125 these examples. This would allow it to mimic the behavior of a better model, without having to pay  
126 for each query.

---

<sup>1</sup>In the sense of the *Chain of Thought* approach, as LLM's actually *thinking* is a topic of debate.

<sup>2</sup>*Easy-to-consider* improvements such as prompting the model indefinitely until it generates a repaired program that passes all tests or having direct access to the problem data's expected solution are not considered in this section, since they might not reflect real-life scenarios with a fixed budget, which makes it infeasible to prompt the model indefinitely.

---

127 **Acknowledgements**

128 Week 10's slides and listed references.

129 **References**

130 [1] Nachiket Kotalwar, Alkis Gotovos, and Adish Singla. Hints-in-browser: Benchmarking language  
131 models for programming feedback generation, 2024.

Following is the setup of a problem in Python. It contains the description and a sample testcase.

```
[Problem Starts]
{problem_data}
[Problem Ends]
```

Following is the student's buggy code for this problem:

```
[Buggy Code Starts]
{buggy_program}
[Buggy Code Ends]
```

Previously, a Large Language Model was asked to review the student's buggy code and generate a repaired program based on the problem description. The repaired program is as follows:

```
[Repaired Code Starts]
{repaired_program}
[Repaired Code Ends]
```

Nevertheless, this repaired program is not guaranteed to be correct. For you to review whether this repaired program is correct or not, the results of the test cases for this problem are provided below:

```
[Testcases results of repaired code Starts]
{testcases_results}
[Testcases results of repaired code Ends]
```

Based on all of the above information, create an explanation written inside placeholders [EXP] and [/EXP] that analyzes the student's buggy code, the provided repaired program and the testcases results, to come up step-by-step with the best approach to provide a hint to the student. After that, provide a concise single-sentence hint to the student about one bug in the student's buggy code, which allows him to get closer to the correct program, but without directly giving him the code or detailed explanations. Output your hint between [HINT] and [/HINT]. Make sure to always use the [HINT] and [/HINT] placeholders with a hint inside. Do not output anything else inside the [HINT] placeholders that is not part of your chosen hint. Do not refer to the student as "the student", but with "you" if strictly necessary. Always start your hints with "Consider" or "Think about".

Figure 1: Advanced prompt assigned to the variable `user_message_nus_hint_advanced` in `project_part1_prompts.py`, which is used in (I.7) and (I.8). Notice how CoT prompting is being used by constraining the model to generate a Chain of Thought inside the [EXP] placeholders. Moreover, in order to constrain the kind of hints it should give, the model is constrained to always start its hints with "Consider" or "Think about".

```
[Problem Starts]
Duplicate elimination - Write a function remove_extras(lst) that...
[Problem Ends]
```

```
[Buggy Code Starts]
def remove_extras(lst):
    new = []
    for x in lst:
        if lst.count(x) > 1:
            new += []
        else:
            new = new.append(x)
    return new
[Buggy Code Ends]
```

Previously, a Large Language Model was asked to review the student's buggy code and generate a repaired program based on the problem description. The repaired program is as follows:

```
[Repaired Code Starts]
def remove_extras(lst):
    new = []
    for x in lst:
        if x not in new:
            new.append(x)
    return new
[Repaired Code Ends]
```

Nevertheless, this repaired program is not guaranteed to be correct. For you to review whether this repaired program is correct or not, the results of the test cases for this problem are provided below:

```
[Testcases results of repaired code Starts]
* 6 out of 6 testcases passed.
[Testcases results of repaired code Ends]
```

Based on all of the above information, create an explanation written inside placeholders [EXP] and [/EXP] that analyzes the student's buggy code, the provided repaired program and the testcases results, to come up step-by-step with the best approach to provide a hint to the student. After that, provide a concise single-sentence hint to the student about one bug in the student's buggy code, which allows him to get closer to the correct program, but without directly giving him the code or detailed explanations. Output your hint between [HINT] and [/HINT]. Make sure to always use the [HINT] and [/HINT] placeholders with a hint inside. Do not output anything else inside the [HINT] placeholders that is not part of your chosen hint. Do not refer to the student as "the student", but with "you" if strictly necessary. Always start your hints with "Consider" or "Think about".

Figure 2: An example of the advanced prompt being used to query GPT-4o-mini. Notice that the repaired program created by GPT-4o-mini passes 6 out of 6 test cases, thus it's a correct repaired program that can be used as context to generate a hint. *Prompt was shortened for brevity.*

[EXP]The student's buggy code attempts to remove duplicates but has a couple of fundamental issues. First, it uses `lst.count(x) > 1` to check if an element is duplicated, which is not necessary for the goal of this function. This approach can be inefficient because `count` iterates through the list every time it is called, leading to a time complexity of  $O(n^2)$  in the worst case. Additionally, in the line `new = new.append(x)`, `append` returns `None`, so `new` ends up being `None` instead of the expected list, thereby causing an error in the function.

The repaired code correctly uses a mechanism to build a new list by checking if each element is already in the new list and only appending it if it's not present. Since all provided test cases for the repaired code passed, it indicates that this approach correctly meets the problem's requirements.

To help you create a functioning solution, focus on ensuring that your program checks for the existence of elements in the new list appropriately and maintains the order of the first occurrences without excessive iterations.

[/EXP]

[HINT]Consider how you are currently checking for duplicates in your code and whether this aligns with the goal of creating a new list that preserves order while eliminating extras.

[/HINT]

Figure 3: Response given by GPT-4o-mini to the advanced prompt in Figure 2.

```
[Problem Starts]
Duplicate elimination - Write a function remove_extras(lst) that...
[Problem Ends]
```

```
[Buggy Code Starts]
def remove_extras(lst):
    new = []
    for x in lst:
        if lst.count(x) > 1:
            new += []
        else:
            new = new.append(x)
    return new
[Buggy Code Ends]
```

Previously, a Large Language Model was asked to review the student's buggy code and generate a repaired program based on the problem description. The repaired program is as follows:

```
[Repaired Code Starts]
def remove_extras(lst):
    new = []
    for x in lst:
        if lst.count(x) == 1:
            new.append(x)
    return new
[Repaired Code Ends]
```

Nevertheless, this repaired program is not guaranteed to be correct. For you to review whether this repaired program is correct or not, the results of the test cases for this problem are provided below:

```
[Testcases results of repaired code Starts]
* 1 out of 6 testcases passed.
[Testcases results of repaired code Ends]
```

Based on all of the above information, create an explanation written inside placeholders [EXP] and [/EXP] that analyzes the student's buggy code, the provided repaired program and the testcases results, to come up step-by-step with the best approach to provide a hint to the student. After that, provide a concise single-sentence hint to the student about one bug in the student's buggy code, which allows him to get closer to the correct program, but without directly giving him the code or detailed explanations. Output your hint between [HINT] and [/HINT]. Make sure to always use the [HINT] and [/HINT] placeholders with a hint inside. Do not output anything else inside the [HINT] placeholders that is not part of your chosen hint. Do not refer to the student as "the student", but with "you" if strictly necessary. Always start your hints with "Consider" or "Think about".

Figure 4: An example of the advanced prompt being used to query Phi-3-mini. Notice that the repaired program created by Phi-3-mini passes 1 out of 6 test cases, thus it's not a completely correct repaired program. *Prompt was shortened for brevity.*

[EXP] The student's buggy code attempts to remove duplicates by checking if the count of an element in the list is greater than 1. However, the use of 'new += []' in the if condition is incorrect because it tries to concatenate an empty list to 'new', which does not affect the list at all. Moreover, the else clause should not reassign 'new' to the result of 'new.append(x)', as this would overwrite the list with a single element, which is not the intended behavior. The repaired code correctly appends unique elements to 'new', but it does not preserve the relative order of the elements, which is a requirement of the problem statement.

[/EXP]

[EXP] The test cases results indicate that the repaired program does not pass all the test cases, suggesting that there is still an issue with the logic that needs to be addressed. The problem with the repaired code is that it does not maintain the relative order of the elements in the original list, which is a key requirement of the problem statement.

[/EXP]

[HINT] Consider using a data structure that maintains the order of elements as you iterate through the list and check for duplicates.

Figure 5: Response given by Phi-3-mini to the advanced prompt in Figure 4. Notice that Phi-3-mini fails to terminate the hint with the [/HINT] placeholder, even though it was asked for in the prompt. The provided hint however is of good quality.